
ECE 175B - Variational AutoEncoder

Anonymous Author(s)

Affiliation

Address

email

Abstract

Aleksandar Jeremic

This paper summarizes a Variational Auto Encoder with 4 Convolutional layers and fully connected layers to encode and decode, the depth was originally much larger but many issues with exploding and vanishing gradients along with long training time cause a pivot to smaller depth. The VAE was trained on chest xray images of size 128x128 and 256x256.

1 Introduction

In this paper a 4 layer convolutional VAE with a linear layer was used along with a latent space of dimension 2048 to reconstruct chest xray images, the VAE uses LeakyReLU and Batchnorm to help with vanishing and exploding gradients which was a major concern in the beginning of this work when there were over 8 layers in encode and decode stages

2 Hyper Parameters

In this paper I used both 256x256 and 128x128 scaled images from the chest xray dataset, Stochastic Gradient Descent and Adam optimization (Adam was the best) was used with a learning rate of 0.001 a batch size of 32. The VAE was implemented with multiple different epochs and the best results were chosen, the number of epochs of training ranged from 10-250 for the 128x128 simplistic model. The loss function involved binary cross entropy loss along with KL Divergence with a weight of 0.001, I also changed the size of the latent dim and found around 200 to give the best results.

3 Training loss Curves

We can see a training loss curve stabilizing after about 10 epochs, in this run the training loss began at a very small value and this iteration had problems with exploding and vanishing gradients which then destroy the training loss.

The loss also often diverged due to various reasons such as vanishing and exploding gradients and hyperparameter tuning




Figure 1: training loss

4 Reconstructed Images

Below are the images that were sent into the VAE, from there they are sent back to decode, these reconstructed images are fairly accurate however fail to recognize some of the ribs of the xrays, this is reasonable considering the most notable section of the xray is around the shoulders where there exists lots of differences in the way the shoulders are positioned. This is also quite a powerful reconstruction considering the small dataset which we trained on, only a few thousand images of size 256x256/128x128 and we could actually reconstruct some basic chest xray images!




Figure 2: An example image.

here is the reconstructed images.




Figure 3: An example image.

5 FID score and IS Loss

FID and IS scoring is a form of accuracy for generative networks such as GANs and VAEs, FID measures the difference between real and generated images and compares how similar they are, IS scores are the diversity and quality of the generated images. I used scoring code directly from the website provided <https://pytorch-ignite.ai/blog/gan-evaluation-with-fid-and-is/>

I had alot of trouble with working FID and IS scores, I would consistently get an error of having an imaginary component and would crash the runtime, this would cause me to lose images which made it very hard to gather consistent results when my training would often diverge. I believe this was in part due to the main exploding and vanishing gradients which could be solved using a larger dataset or a shallower algorithm.

6 Conclusion

The VAE is a Machine Learning Model which uses a latent space representation of input images to understand information about the image, from the latent space the trained model can recreate images similar to the ones that were inputted and trained by the model. Alternatively sampling from a distribution in the latent space is equivalent to picking a latent space representation of the training set and then reconstruction from that sample is generating a similar image to the training set.

7 Code

```

51 import torch
52 import torch.nn as nn
53 import torch.optim as optim
54 from torch.utils.data import DataLoader
55 from torchvision import transforms
56 from torchvision.datasets import ImageFolder
57 from torchvision.utils import save_image
58
59 import torch.nn.functional as F
60
61
62
63 from google.colab import drive
64 drive.mount('/content/drive')
65
66
67
68 #!unzip /content/drive/MyDrive/chest_xray.zip -d /content/drive/
69 MyDrive/
70
71 import os
72
73
74 dataset_dir = '/content/drive/MyDrive/chest_xray/chest_xray'
75
76 data_dir = '/content/drive/MyDrive/chest_xray'
77

```

```

78 base_path = '/content/drive/MyDrive/chest_xray'
79
80 data_transforms = transforms.Compose([
81     transforms.Grayscale(),
82     transforms.Resize((128, 128)),
83     transforms.ToTensor(),
84 ])
85
86 train_dir = os.path.join(dataset_dir, 'train')
87 test_dir = os.path.join(dataset_dir, 'test')
88 val_dir = os.path.join(dataset_dir, 'val')
89
90 USE_GPU = True
91 dtype = torch.float32
92
93 if USE_GPU and torch.cuda.is_available():
94     device = torch.device('cuda')
95 else:
96     device = torch.device('cpu')
97
98 print_every = 100
99
100 print('using device:', device)
101
102 train_dataset = ImageFolder(train_dir, transform=data_transforms)
103 test_dataset = ImageFolder(test_dir, transform=data_transforms)
104 val_dataset = ImageFolder(val_dir, transform=data_transforms)
105
106 batch_size = 32
107 train_loader = DataLoader(train_dataset, batch_size=batch_size,
108     shuffle=True)
109 test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=
110     False)
111 val_loader = DataLoader(val_dataset, batch_size=batch_size, shuffle=
112     False)
113
114 class VAE(nn.Module):
115     def __init__(self):
116         super(VAE, self).__init__()
117
118         self.conv1_1 = nn.Conv2d(1, 64, (3, 3), padding=1, stride=2)
119         self.conv1_2 = nn.Conv2d(64, 64, (3, 3), padding=1)
120
121         self.conv2_1 = nn.Conv2d(64, 128, (3, 3), padding=1, stride=2)
122         self.conv2_2 = nn.Conv2d(128, 128, (3, 3), padding=1)
123
124         self.conv3_1 = nn.Conv2d(128, 256, (3, 3), padding=1, stride
125             =2)
126         self.conv3_2 = nn.Conv2d(256, 256, (3, 3), padding=1)
127         self.conv3_3 = nn.Conv2d(256, 256, (3, 3), padding=1)
128
129         self.conv4_1 = nn.Conv2d(256, 256, (3, 3), padding=1, stride
130             =2)
131         self.conv4_2 = nn.Conv2d(256, 256, (3, 3), padding=1)
132         self.conv4_3 = nn.Conv2d(256, 256, (3, 3), padding=1)
133
134         self.conv5_1 = nn.Conv2d(512, 512, (3, 3), padding=1, stride
135             =2)
136         self.conv5_2 = nn.Conv2d(512, 512, (3, 3), padding=1)
137         self.conv5_3 = nn.Conv2d(512, 512, (3, 3), padding=1)
138         self.conv5_4 = nn.Conv2d(512, 512, (3, 3), padding=1)
139
140         self.bn1_1 = nn.BatchNorm2d(64)
141         self.bn1_2 = nn.BatchNorm2d(64)
142         self.bn2_1 = nn.BatchNorm2d(128)

```

```

143         self.bn2_2 = nn.BatchNorm2d(128)
144
145         self.bn3_1 = nn.BatchNorm2d(256)
146         self.bn3_2 = nn.BatchNorm2d(256)
147         self.bn3_3 = nn.BatchNorm2d(256)
148
149         self.bn4_1 = nn.BatchNorm2d(256)
150         self.bn4_2 = nn.BatchNorm2d(256)
151         self.bn4_3 = nn.BatchNorm2d(256)
152
153         self.bn5_1 = nn.BatchNorm2d(512)
154         self.bn5_2 = nn.BatchNorm2d(512)
155         self.bn5_3 = nn.BatchNorm2d(512)
156         self.bn5_4 = nn.BatchNorm2d(512)
157
158         self.relu = nn.LeakyReLU(0.01)
159         self.sigmoid = nn.Sigmoid()
160
161         self.fc1 = nn.Linear(256 * 8 * 8, 512)
162         self.fc2 = nn.Linear(512, 256 * 8 * 8)
163
164         self.fc3 = nn.Linear(256 * 8 * 8, 512)
165         self.fc4 = nn.Linear(256 * 8 * 8, 512)
166
167         self.deconv1_1 = nn.ConvTranspose2d(512, 512, (3, 3), padding
168         =1)
169         self.deconv1_2 = nn.ConvTranspose2d(512, 512, (3, 3), padding
170         =1)
171         self.deconv1_3 = nn.ConvTranspose2d(512, 512, (3, 3), padding
172         =1)
173         self.deconv1_4 = nn.ConvTranspose2d(512, 512, (3, 3), padding
174         =1, stride=2, output_padding=1)
175
176         self.deconv2_1 = nn.ConvTranspose2d(256, 256, (3, 3), padding
177         =1)
178         self.deconv2_2 = nn.ConvTranspose2d(256, 256, (3, 3), padding
179         =1)
180         self.deconv2_3 = nn.ConvTranspose2d(256, 256, (3, 3), padding
181         =1, stride=2, output_padding=1)
182
183         self.deconv3_1 = nn.ConvTranspose2d(256, 256, (3, 3), padding
184         =1)
185         self.deconv3_2 = nn.ConvTranspose2d(256, 256, (3, 3), padding
186         =1)
187         self.deconv3_3 = nn.ConvTranspose2d(256, 128, (3, 3), padding
188         =1, stride=2, output_padding=1)
189
190         self.deconv4_1 = nn.ConvTranspose2d(128, 128, (3, 3), padding
191         =1)
192         self.deconv4_2 = nn.ConvTranspose2d(128, 64, (3, 3), padding
193         =1, stride=2, output_padding=1)
194
195         self.deconv5_1 = nn.ConvTranspose2d(64, 64, (3, 3), padding=1)
196         self.deconv5_2 = nn.ConvTranspose2d(64, 1, (3, 3), padding=1,
197         stride=2, output_padding=1)
198
199         self.bn1_1_deconv = nn.BatchNorm2d(512)
200         self.bn1_2_deconv = nn.BatchNorm2d(512)
201         self.bn1_3_deconv = nn.BatchNorm2d(512)
202         self.bn1_4_deconv = nn.BatchNorm2d(512)
203
204         self.bn2_1_deconv = nn.BatchNorm2d(256)
205         self.bn2_2_deconv = nn.BatchNorm2d(256)
206         self.bn2_3_deconv = nn.BatchNorm2d(256)
207

```

```

208         self.bn3_1_deconv = nn.BatchNorm2d(256)
209         self.bn3_2_deconv = nn.BatchNorm2d(256)
210         self.bn3_3_deconv = nn.BatchNorm2d(128)
211
212         self.bn4_1_deconv = nn.BatchNorm2d(128)
213         self.bn4_2_deconv = nn.BatchNorm2d(64)
214
215         self.bn5_1_deconv = nn.BatchNorm2d(64)
216         self.bn5_2_deconv = nn.BatchNorm2d(1)
217
218     def encode(self, x):
219         x = self.conv1_1(x)
220         x = self.bn1_1(x)
221         x = self.relu(x)
222
223         x = self.conv2_1(x)
224         x = self.bn2_1(x)
225         x = self.relu(x)
226
227         x = self.conv3_1(x)
228         x = self.bn3_1(x)
229         x = self.relu(x)
230
231         x = self.conv4_1(x)
232         x = self.bn4_1(x)
233         x = self.relu(x)
234
235         x = x.view(x.size(0), -1)
236         mean = self.fc3(x)
237         log_var = self.fc4(x)
238         return mean, log_var
239
240     def reparameterize(self, mu, logvar):
241         std = torch.exp(0.5 * logvar)
242         eps = torch.randn_like(std)
243         return mu + eps * std
244
245     def decode(self, z):
246         x = self.fc2(z)
247         x = x.view(x.size(0), 256, 8, 8)
248
249         x = self.deconv2_3(x)
250         x = self.bn2_3_deconv(x)
251         x = self.relu(x)
252
253         x = self.deconv3_3(x)
254         x = self.bn3_3_deconv(x)
255         x = self.relu(x)
256
257         x = self.deconv4_2(x)
258         x = self.bn4_2_deconv(x)
259         x = self.relu(x)
260
261         x = self.deconv5_2(x)
262         x = self.bn5_2_deconv(x)
263         x = self.relu(x)
264
265         x = self.sigmoid(x)
266         return x
267
268     def forward(self, x):
269         mu, log_var = self.encode(x)
270         z = self.reparameterize(mu, log_var)
271         x_reconstructed = self.decode(z)
272         return x_reconstructed, mu, log_var

```

```

273 model = VAE().to(device)
274
275
276 optimizer = optim.Adam(model.parameters(), lr=0.001)
277
278 kld_rate = .001
279
280
281
282 num_epochs = 20
283
284 for epoch in range(num_epochs):
285     running_loss = 0.0
286     for batch_idx, (data, _) in enumerate(train_loader):
287         data = data.to(device)
288
289         x, mu, log_var = model(data)
290
291         loss = F.binary_cross_entropy(x, data) - 0.5 * kld_rate * torch.
292             sum(1 + log_var - mu.pow(2) - log_var.exp())
293
294         optimizer.zero_grad()
295         loss.backward()
296         optimizer.step()
297
298         running_loss += loss.item()
299
300         if batch_idx % 100 == 99:
301             print('Epoch [{}/{}], Step [{}/{}], Loss: {:.4f}'
302                 .format(epoch + 1, num_epochs, batch_idx + 1, len(
303                     train_loader), running_loss / 100))
304             running_loss = 0.0
305
306
307 torch.save(model.state_dict(), "vae_model.pth")
308
309 torch.save(model.state_dict(), os.path.join(os.getcwd(), "model2.pth")
310 )
311 model.load_state_dict(torch.load(os.path.join(os.getcwd(), "model2.pth
312 ")))
313
314 with torch.no_grad():
315     model.eval()
316     for i, (data, _) in enumerate(test_loader):
317         data = data.to(device)
318         reconstructed, _, _ = model(data)
319         if i == 0:
320             save_image(data, os.path.join(data_dir, "original_images.
321                 png"))
322             save_image(reconstructed, os.path.join(data_dir, "
323                 reconstructed_images.png"))
324         break
325

```